Prompt for Conversation completeness :

Implement Conversation Completeness Metric

Objective

Add a new RAG evaluation metric named "Conversation Completeness" using the same architecture already used by the existing metrics.

Maintain UI consistency with all existing evaluation metrics.

Do not redesign existing components.

Metric Name

Conversation Completeness

Metric Purpose

Conversation Completeness evaluates whether a multi-turn conversation fully addresses the user's needs, questions, and intents by the end, with no unanswered questions, ignored requests, or unresolved follow-ups.

Teaching Mode

Overview

Explain:

- What conversation completeness means in AI systems

- Why incomplete conversations frustrate users

- Why enterprises measure whether user goals were fully resolved

Advantages

- Detects unanswered or ignored user requests

- Improves user satisfaction and resolution rates

- Measures the whole conversation, not just single replies

Limitations

- User intent can be ambiguous or change mid-conversation
- Some conversations are legitimately open-ended
- Judging "complete" can be subjective

Industry Usage

- Customer support
- Banking assistants
- Healthcare assistants
- Educational systems
- Hiring and HR helpdesks

Real World Applications

- Support chatbots resolving tickets
- Banking virtual assistants
- Student tutoring chatbots
- HR copilots answering employee queries

Best Practices

- Extract every user intent from the conversation
- Check each intent for a clear, satisfying answer
- Test with multi-turn, multi-question scenarios
- Review unresolved conversations regularly

Common Mistakes

- Evaluating only the last reply
- Ignoring questions asked earlier in the conversation
- Treating a long conversation as a complete one

Good Example

Conversation:

User: What are your support hours, and how do I reset my password?

Assistant: Support is available 9 AM to 6 PM, Monday to Friday. To reset your password, click "Forgot Password" on the login page and follow the email link.

Result: High Completeness

Bad Example

Conversation:

User: What are your support hours, and how do I reset my password?

Assistant: Support is available 9 AM to 6 PM, Monday to Friday.

Result: Low Completeness (password question ignored)

Score Interpretation

0.90 - 1.00 Fully Complete

0.75 - 0.89 Mostly Complete

0.50 - 0.74 Partially Complete

Below 0.50 Largely Incomplete

Sample Library

Create Good and Bad multi-turn examples for:

- Hiring
- Education
- Banking
- Customer Support
- Healthcare

Q&A Prep

Add beginner-friendly questions:

- What is conversation completeness?

- Why is it different from answer relevancy?

- Can a conversation be accurate but incomplete?

- How do enterprises improve resolution rates?

UI Requirements

Reuse:

- Existing metric selector

- Existing teaching mode

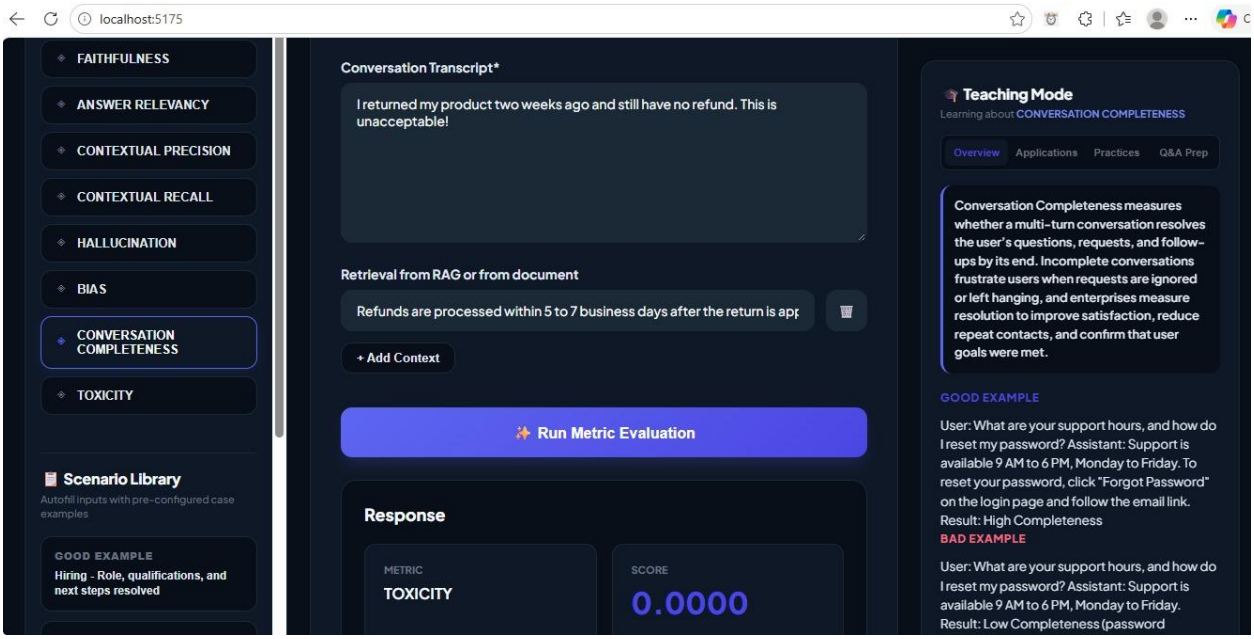
- Existing sample library

- Existing evaluation workflow

- Existing result cards

If the existing workflow only accepts a single question/answer, extend the input minimally to support multi-turn conversation text, without redesigning it.

Ensure the Conversation Completeness metric feels identical to the existing metrics already implemented.



Prompt for Toxicity :

Implement Toxicity Metric

Objective

Add a new RAG evaluation metric named "Toxicity" using the same architecture already used by the existing metrics.

Maintain UI consistency with all existing evaluation metrics.

Do not redesign existing components.

Metric Name

Toxicity

Metric Purpose

Toxicity evaluates whether the generated answer contains harmful, offensive, abusive, hateful, insulting, or threatening language.

Teaching Mode

Overview

Explain:

- What toxicity means in AI systems

- Why toxic output is dangerous

- Why enterprises screen model responses for harmful language

Advantages

- Protects users from harmful content

- Protects brand reputation

- Supports safe and responsible AI deployment

Limitations

- Sarcasm and slang are hard to judge

- Cultural and contextual differences affect interpretation

- Quoting or discussing toxic content is not the same as being toxic

Industry Usage

- Customer support

- Social platforms and community moderation

- Educational systems

- Healthcare assistants

- Banking and HR helpdesks

Real World Applications

- Support chatbots

- Moderation assistants

- Student-facing tutoring bots

- Employee-facing HR copilots

Best Practices

- Test with adversarial and provocative prompts

- Combine automated scoring with human review

- Set clear content policies and thresholds

- Monitor outputs continuously after release

Common Mistakes

- Checking only for obvious profanity

- Ignoring subtle insults or condescension

- Testing only polite user inputs

Good Example

Question:

My order is late again. This is ridiculous!

Generated Answer:

I'm sorry for the delay. Let me check your order status and fix this right away.

Result:

Low Toxicity

Bad Example

Question:

My order is late again. This is ridiculous!

Generated Answer:

Stop complaining. Maybe if you weren't so impatient, you wouldn't have problems.

Result:

High Toxicity

Score Interpretation (higher = safer)

0.90 - 1.00 Safe / Non-Toxic

0.75 - 0.89 Mostly Safe

0.50 - 0.74 Potentially Toxic

Below 0.50 Significant Toxicity Risk

Sample Library

Create Good and Bad examples for:

- Hiring

- Education

- Banking

- Customer Support

- Healthcare

Q&A Prep

Add beginner-friendly questions:

- What is AI toxicity?

- Why is toxic output dangerous?

- Can a helpful answer still be toxic?

- How do enterprises reduce toxic responses?

UI Requirements

Reuse:

- Existing metric selector

- Existing teaching mode

- Existing sample library

- Existing evaluation workflow

- Existing result cards

Ensure the Toxicity metric feels identical to the existing metrics already implemented.